

# Fitting mathematical models to data using Non-linear Least-Squares

Samraat Pawar

*Department of Life Sciences*  
*(Silwood Park)*

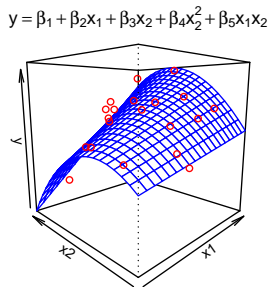
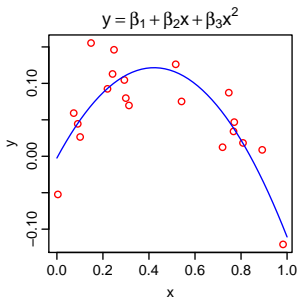
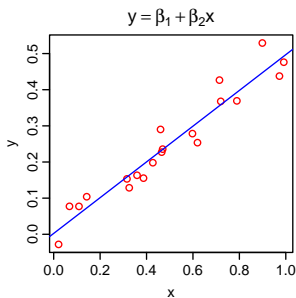
**Imperial College**  
**London**

January 20, 2018

# OUTLINE

- Why Non-Linear Least Squares regression / fitting?
- The NLLS fitting method
- NLLS in R
- Afternoon practicals overview

# LINEAR MODELS ARE GREAT



- These are *all* good *Linear Models* (really?!)
- The data can be modelled (aka "fitted to a mathematical model") as a *linear combination of variables and coefficients*
- Easily fitted using *Ordinary Least Squares* (OLS) regression
- Linear models can *include curved responses* (e.g. polynomial regression)
- OK, so then *why Non-Linear Least Squares (NLLS) fitting?*

# WHY NLLS? – FIRST, WHAT MAKES A MODEL NON-LINEAR?

- OLS can be used to fit both linear and nonlinear *equations* that *intrinsically linear*, e.g.,
  - Straight line:  $y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$
  - Polynomial:  $y_i = \exp(\beta_0) + \beta_1 x_i + \beta_2 x_i^2 + \varepsilon_i$
- Indeed, for OLS to work, we need *intrinsic linearity* — i.e., the equation to be fitted (model) should be *linear in the parameters*
- Are these models linear in their *parameters*?
  - $y_i = \beta_0 + \beta_1 x_i^{\beta_2} + \varepsilon_i$
  - $y_i = \beta_0 e^{\beta_2 x_i} + \varepsilon_i$

NO!

# SO WHAT — WHY IS INTRINSIC NON-LINEARITY A PROBLEM?

## Recall what the Least Squares method does:

- Consider a predictor  $x$ , data  $y$ ,  $n$  observations, and a model that we want to fit to the data:

$$f(x_i, \beta) + \varepsilon_i$$

where  $\beta = (\beta_1, \beta_2, \dots, \beta_k)$  are the model's  $k$  parameters


- The objective is to find estimates of values of the  $k$  parameters ( $\hat{\beta}_j$ ) that minimize the sum ( $S$ ) of squared residuals ( $r_i$ ) (AKA RSS):

$$S = \sum_{i=1}^n [y_i - f(x_i, \beta)]^2 = \sum_{i=1}^n r_i^2$$

# THE LEAST-SQUARES SOLUTION

OLS minimizes the *sum* of the *squared* residuals

# IF THE MODEL IS LINEAR, THE SOLUTION IS EASY USING ALGEBRA



$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$$

$$9.50 = 5 + 4 \times 1 + 0.50$$

$$11.00 = 5 + 4 \times 2 - 2.00$$

$$19.58 = 5 + 4 \times 3 + 2.58$$

$$20.00 = 5 + 4 \times 4 - 1.00$$

$$\beta_0 = 5; \beta_1 = 4$$

# INTRINSIC NON-LINEARITY DOES NOT ALLOW A ALGEBRAIC SOLUTION

- So, then, in an intrinsically non-linear model such as  $y_i = \beta_0 e^{\beta_2 x_i} + \varepsilon_i$  the derivatives  $\frac{\partial r_i}{\partial \beta_j}$  are naughty
- That is, they are functions of both  $x$  and the parameters  $\beta_j$ , so the gradient equations do not have a solution like the OLS case
- So the nice trick of solving  $\mathbf{Y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$  is impossible *mathematically*



## SO — ENTER NLLS

But we can use brute-force computation to find close-to-optimal least squares minimization!

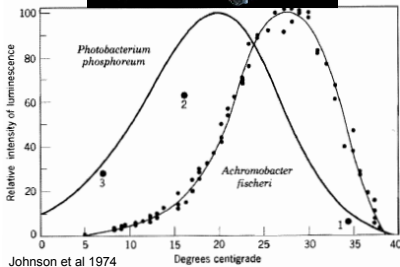
- Choose initial values for the parameters we want to estimate ( $\beta_j$ 's)
- Then, “refine” the parameters *iteratively* by calculating  $\frac{\partial r_i}{\partial \beta_j}$  *approximately* — this approximation is the *Jacobian* (the gradient), which is a matrix of the  $\frac{\partial r_i}{\partial \beta_j}$ 's
- Whether a refinement has taken place in any step of the iteration is determined by re-calculating the residuals at that step
- Eventually, if it all goes well, we find a combination of  $\beta_j$ 's that is *very close* to the desired solution  $\frac{\partial S}{\partial \beta_j} = 0, j = 0, 1, 2, \dots, k$



# OK, FINE, WHY WOULD I EVER NEED NLLS?

- Many observations in biology are just *not* well-fitted by a linear model
- That is, the underlying biological phenomena/phenomenon are not well-described by a linear equation
- Examples:
  - Logistic population growth
  - Allometric growth
  - Michaelis-Menten biochemical kinetics (two parameters  $V_{\max}$  and  $K_m$ :  $v = \frac{V_{\max}[S]}{K_m + [S]}$ )
  - Responses of metabolic rates to changing temperature
  - Consumer-Resource (e.g., predator-prey) functional responses
  - Time-series data (e.g., fitting a sinusoidal function)
- *Can you think of some examples?*

# EXAMPLE: TEMPERATURE AND METABOLISM



$$B = B_0 \boxed{e^{-\frac{E}{kT}}} f(T, T_{pk}, E_D)$$

$T$  = temperature (K)

$k$  = Boltzmann constant ( $\text{eV K}^{-1}$ )

$E$  = Activation energy (eV)

$T_{pk}$  = Temperature of peak performance

$E_D$  = Deactivation energy (eV)

(J H vant Hoff 1884, S Arrhenius 1889)

- Surely there is more to thermal responses?
  - Oxygen limitation
  - Complexity of metabolic network
  - Hormonal regulation
- *What about alternative models?*

## EXAMPLE: FUNCTIONAL RESPONSES


$$f(x_R) = \frac{ax_R^{q+1}}{1+hax_R^{q+1}} \text{ (Holling, 1959)}$$

$x_R$  = Resource density (Mass / Area or Volume)

$a$  = Search rate (Area or Volume / Time )

$h$  = Handling time

$q$  = Shape parameter (dimensionless)



Note that:

- NLLS fitting can yield  $h < 0$ ,  $q < 0$ , or both
- $h < 0$  is biologically impossible but indicates an upward curving response
- $q < 0$  is biologically unlikely as it indicates a decline in search rate with resource density (but is useful as a measure of deviation away from a type III response)

# NLLS FITTING

So the general procedure is:

- ➊ Start with an initial value for each parameter in the model
- ➋ Generate the curve defined by the initial values
- ➌ Calculate the residual sum-of-squares (rss)
- ➍ Adjust the parameters to make the curve come closer to the data points. This the tricky part — more on this in the next slide
- ➎ Adjust the parameters again so that the curve comes even closer to the points (rss decreases)
- ➏ Repeat 4–5
- ➐ Stop simulations when the adjustments make virtually no difference to the rss

# NLLS FITTING

The *tricky part* — *adjust parameters to make curve come closer to the data points* (step 4) has at least two algorithms:

- The Gauss-Newton algorithm is the default in the `nls` package (part of the `stats` base package) — good in many cases, but doesn't work very well if the model is mathematically weird (the optimization landscape is difficult) and the starting values for parameters are far-off-optimal
- The Levenberg-Marquardt (LM) switches between Gauss-Newton and “gradient descent” and is more robust against starting values that are far-off-optimal — available in R through the `minpack.lm` package  
<http://cran.r-project.org/web/packages/minpack.lm>
- The command is `nlsLM`

# NLLS FITTING

- Once the algorithm has converged (hopefully – but you may be surprised how well it usually works), you need to get the goodness of fit measures
- First, of course, examine the fits visually
- Also, report the best-fit results, including:
  - Sums of deviations of the data points from the final model fit (final RSS)
  - $R^2$
  - Estimated coefficients
  - For each coefficient, standard error (can be used for CI's), t-statistic and corresponding (two-sided) p-value
- The function `summary.nls` will give you all these measures
- Remember, the precise parameter values you obtain will depend in part on the initial values chosen and the convergence criteria
- You may also want to compare multiple models...



# NLLS ASSUMPTIONS

NLLS-regression has all the assumptions of OLS-regression:

- No (in practice, minimal) measurement error in explanatory variable ( $x$ -axis variable)
- Data have constant normal variance — errors in the  $y$ -axis are homogeneously distributed over the  $x$ -axis range
- The measurement/observation error distribution is Gaussian — for example, what would the error distribution of this non-linear model be:  $y_i = \beta_0 e^{\beta_2 x_i} + \varepsilon_i$
- What if the errors are not normal? — use maximum likelihood instead! (e.g., using `nlm` for optimizing/fitting)

# COMPARING MODELS

- It's all about the “Likelihood” of a model:
- That is, the likelihood of a set of parameter values (of a model),  $\theta$ , given outcomes  $x$ , equals the probability of those observed outcomes given those parameter values, that is,

$$\mathcal{L}(\theta|x) = P(x|\theta)$$

# COMPARING MODELS

The easiest thing to do for you is to use information theory (including AIC and BIC) to compare models.

*Both use estimated likelihood of a model*

This is how you can calculate these (using R syntax):

- residuals = Observations - Predictions
- `rss = sum(residuals ^ 2)`
- Then, AIC is  $n * \log((2 * \pi) / n) + n + 2 + n * \log(rss) + 2 * k$   
(what is  $n$  and  $k$ ?)
- And BIC is  $n + n * \log(2 * \pi) + n * \log(rss / n) + (\log(n)) * (k + 1)$
- For both AIC and BIC, If model **A** has AIC lower by 2-3 or more than model **B**, its better — Differences of less than 2-3 don't really matter

Also note that:

- $R^2 = 1 - (rss/tss)$ , where  $tss$  is total sum of squares:  
`tss = sum((Observations - mean(Predictions)) ^ 2)`

# PRACTICALS: INSTRUCTIONS

- All materials are at: <https://goo.gl/8b9vMh>
- You will work with multiple examples and on some Exercises.
- Keep workflow organized in Code, Results, Data as you learnt in the R week!
- You will have demonstrators and me to help you – You may also talk to us about using NLLS for your own project

## MORE NLLS TIPS

- You can use mixed-effects modelling with NLLS in R; the package is `nlme` <https://stat.ethz.ch/R-manual/R-devel/library/nlme/html/nlme.html> (You are probably stuck with the Gauss-Newton algorithm with `nlme` though)
- You can also use Python — look up `lmfit` <https://lmfit.github.io/lmfit-py/index.html>. `python` seems to have a better Levenberg-Marquardt implementation than R

# READINGS AND RESOURCES

- Motulsky, Harvey, and Arthur Christopoulos. Fitting models to biological data using linear and nonlinear regression: a practical guide to curve fitting. OUP USA, 2004.
- Johnson, J. B. & Omland, K. S. 2004 Model selection in ecology and evolution. Trends Ecol. Evol. 19, 101–108.